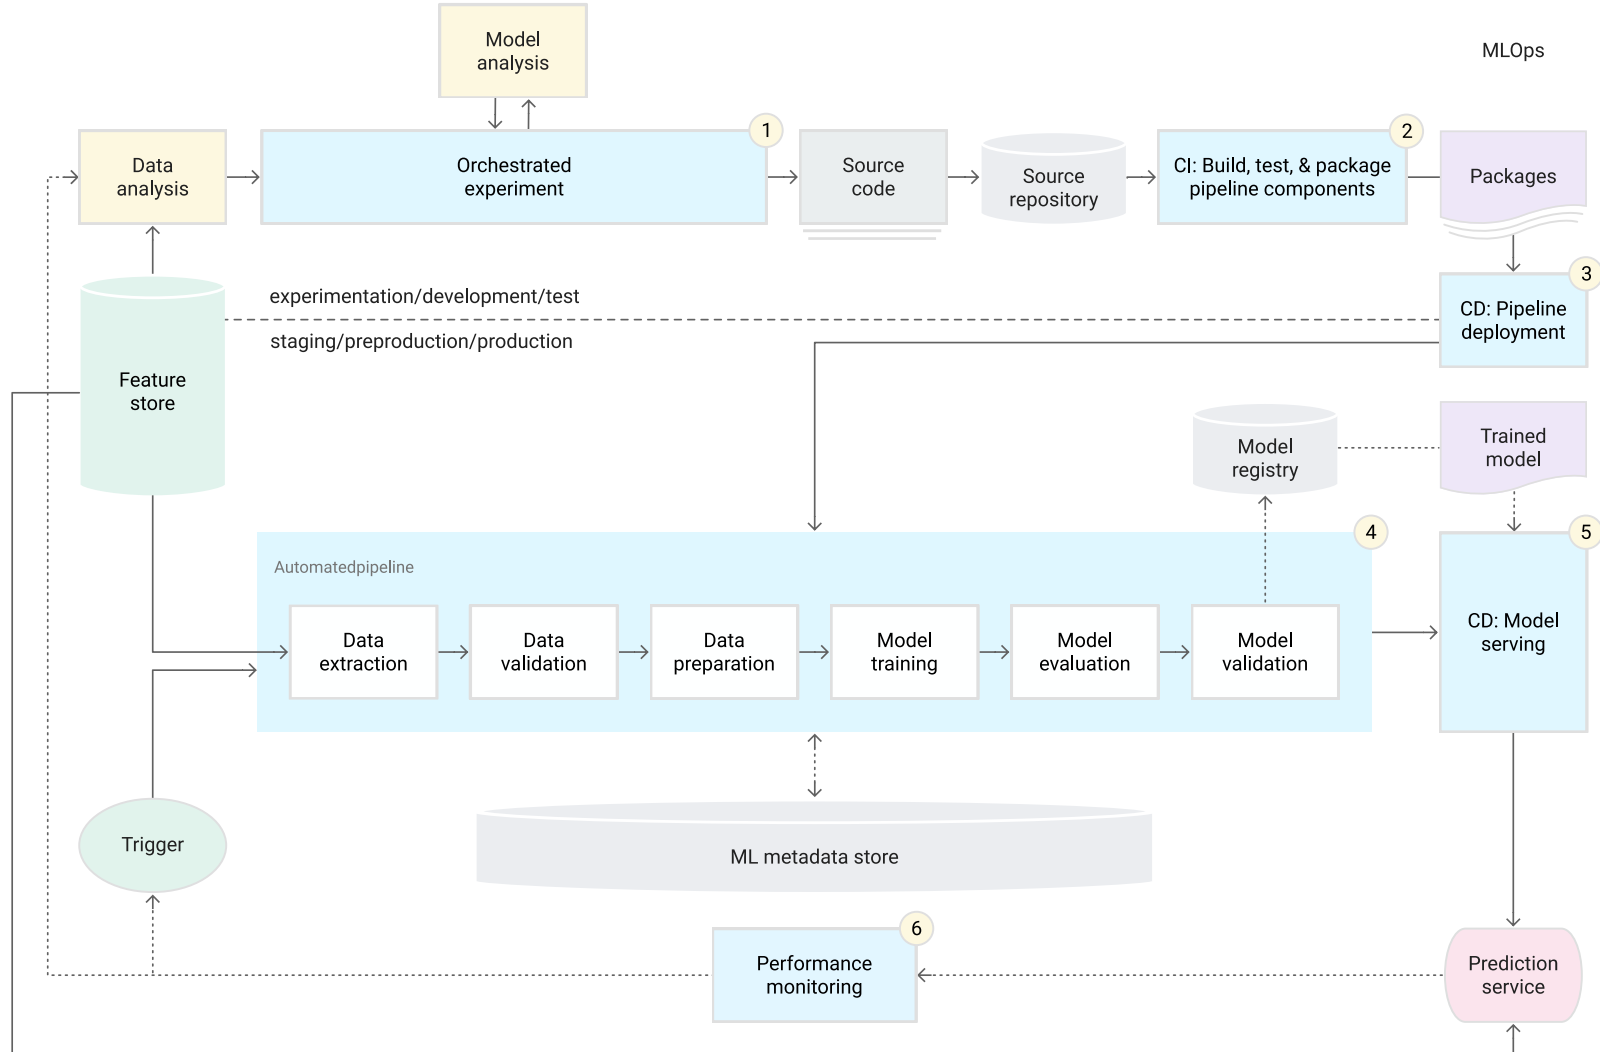


Мониторинг

Семинар 10

Recap MLOps



Ожидания от мониторинга

- Алертинг о сбоях в системе
- Предупреждения о возможных сбоях и проблемах
- Отражение состояния системы/сервера/сервиса/компонента сервиса
- Сбор статистики и визуализаций
- Отчеты
- Дашборды

4 золотых сигнала

SRE handbook от Google:

1. Latency – время, необходимое на обслуживание запроса
2. Traffic – количество запросов, отправляемых в систему
3. Errors – количество ошибок
4. Saturation – насколько полон ваш сервис

The USE

Resource – все компоненты физического/виртуального сервера (CPU, Disk, RAM, etc.)

Utilization – время, затрачиваемое на выполнение задач

Saturation – показатель, указывающий на количество тасок, которые не могут быть выполнены, попадающие при этом в очередь

Errors – количество ошибок

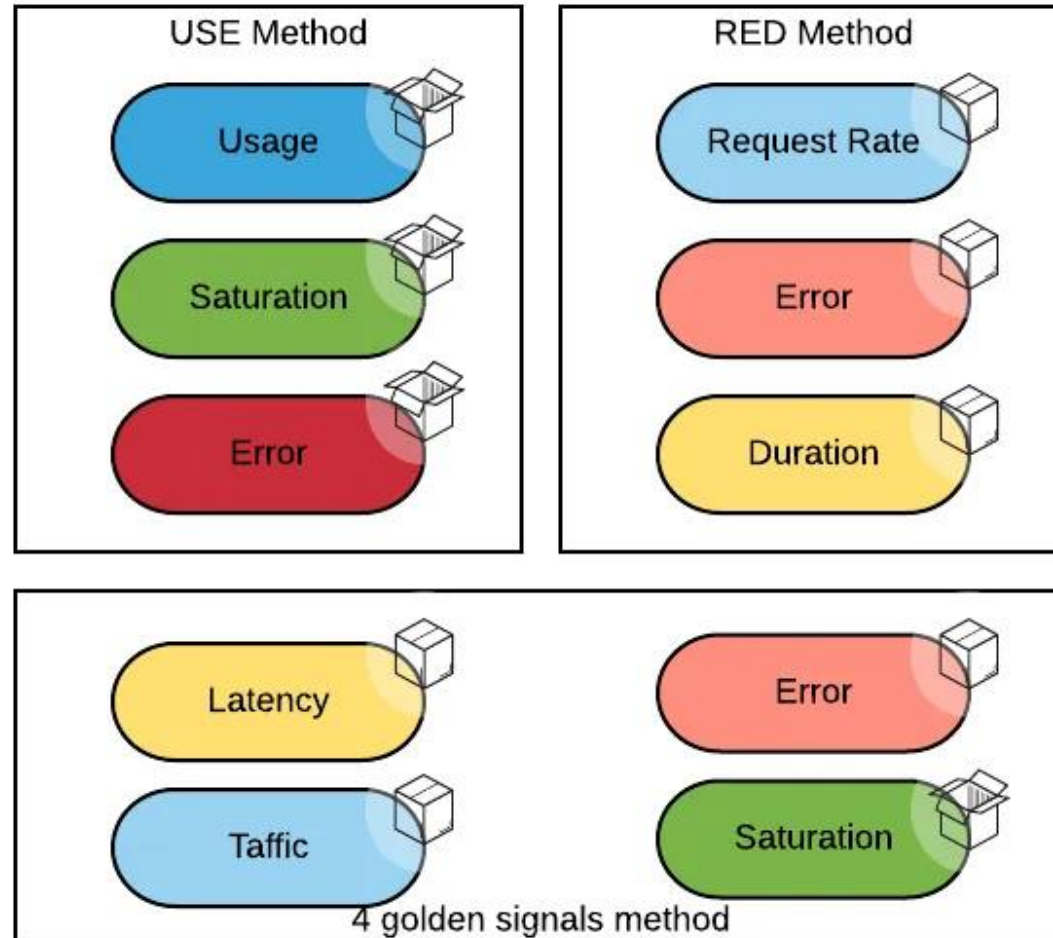
The RED

Rate – количество запросов в секунду

Errors – количество запросов, завершившихся с ошибкой

Duration – количество времени, которое занимает каждый запрос

Все вместе



Метрики

Технические метрики:

Касаются ресурсов (средняя загрузка процессора, занятость диска)

Сервисные метрики:

Касаются запросов (процент успешных запросов, latency)

Бизнес метрики:

MAU, DAU, Clicks, etc

Prometheus

Web-UI

Alertmanager

Prometheus server

Pushgateway



Push & Pull

Push

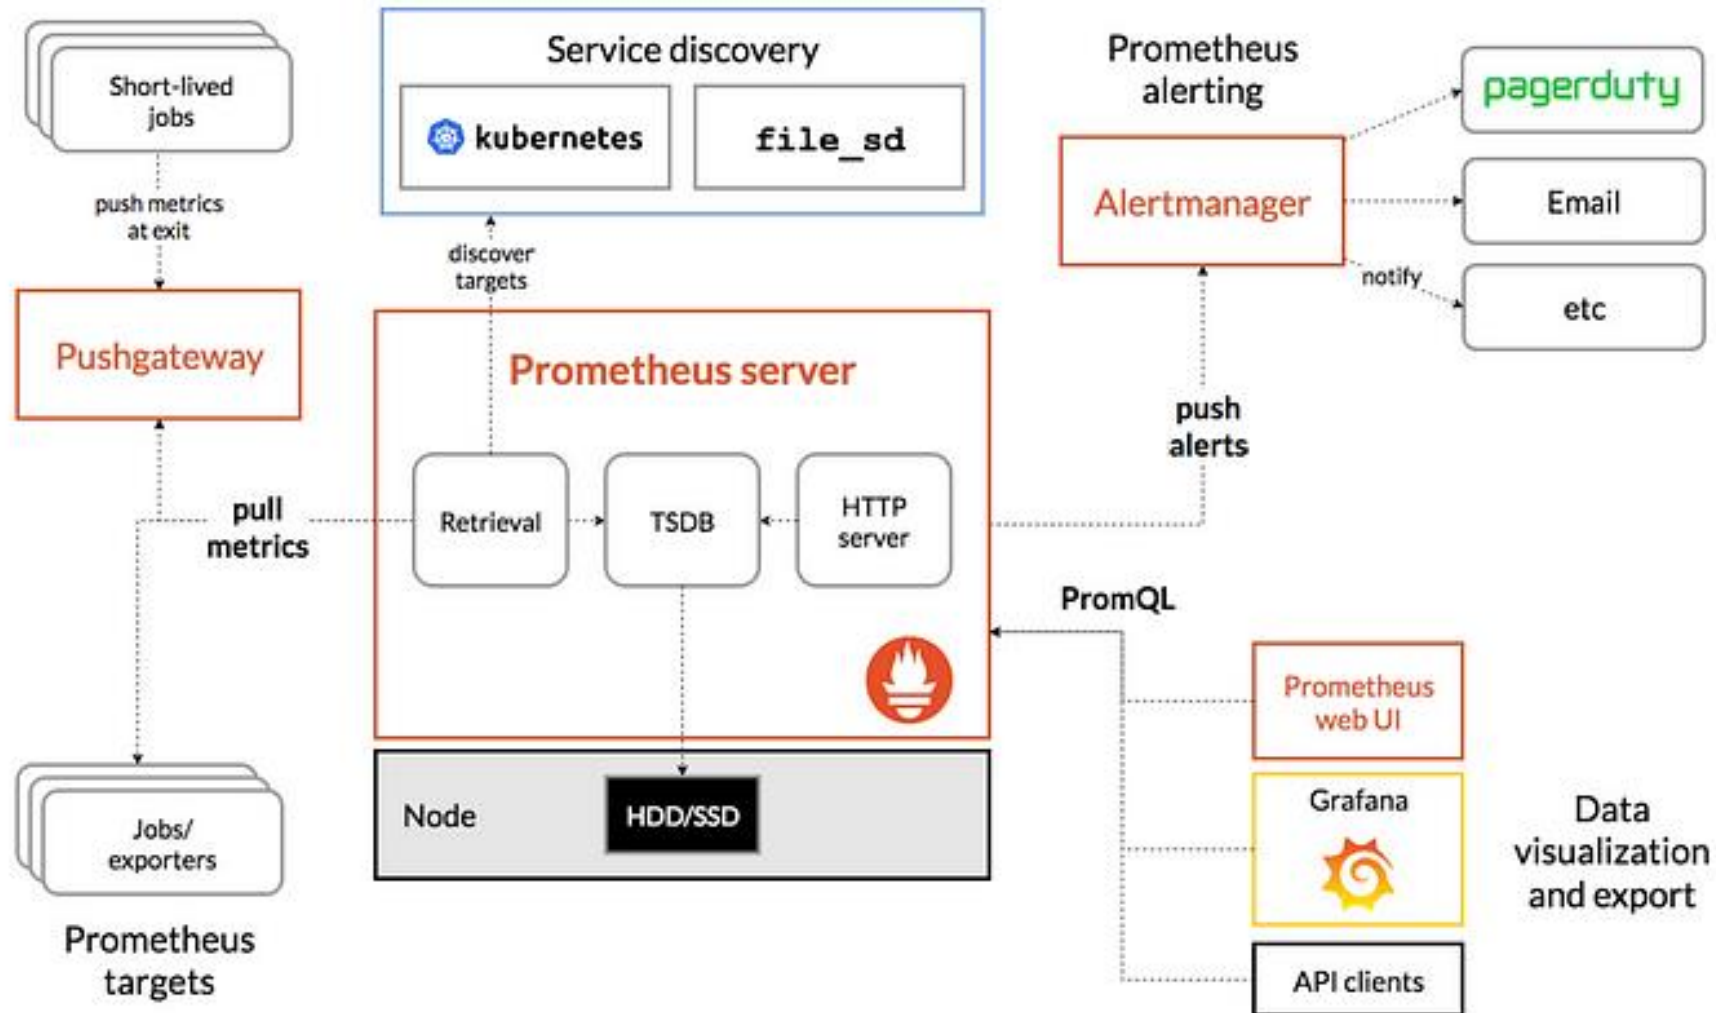
Приложение отправляет
куда-то свои метрики

Pull

Приложение выставляет
endpoint с метриками



Architecture



Типы метрик – Counter

Counter

A *counter* is a cumulative metric that represents a single monotonically increasing counter whose value can only increase or be reset to zero on restart. For example, you can use a counter to represent the number of requests served, tasks completed, or errors.

Do not use a counter to expose a value that can decrease. For example, do not use a counter for the number of currently running processes; instead use a gauge.

Примеры:

- Количество успешных запросов
- Количество ошибок

Типы метрик – Gauge

Gauge

A *gauge* is a metric that represents a single numerical value that can arbitrarily go up and down.

Gauges are typically used for measured values like temperatures or current memory usage, but also "counts" that can go up and down, like the number of concurrent requests.

Примеры:

- Количество машин
- Количество задач в работе

Типы метрик - Histogram

Histogram

A *histogram* samples observations (usually things like request durations or response sizes) and counts them in configurable buckets. It also provides a sum of all observed values.

A histogram with a base metric name of `<basename>` exposes multiple time series during a scrape:

- cumulative counters for the observation buckets, exposed as `<basename>_bucket{le="<upper inclusive bound>"}`
- the **total sum** of all observed values, exposed as `<basename>_sum`
- the **count** of events that have been observed, exposed as `<basename>_count` (identical to `<basename>_bucket{le="+Inf"}` above)

Типы метрик – Summary

Summary

Similar to a *histogram*, a *summary* samples observations (usually things like request durations and response sizes). While it also provides a total count of observations and a sum of all observed values, it calculates configurable quantiles over a sliding time window.

A summary with a base metric name of `<basename>` exposes multiple time series during a scrape:

- streaming **φ -quantiles** ($0 \leq \varphi \leq 1$) of observed events, exposed as `<basename>{quantile="< ϕ >"}`
- the **total sum** of all observed values, exposed as `<basename>_sum`
- the **count** of events that have been observed, exposed as `<basename>_count`

Клиенты

```
from prometheus_client import Counter
c = Counter('my_failures', 'Description of counter')
c.inc()      # Increment by 1
c.inc(1.6)   # Increment by given value
```

```
from prometheus_client import Gauge
g = Gauge('my_inprogress_requests', 'Description of gauge')
g.inc()      # Increment by 1
g.dec(10)    # Decrement by given value
g.set(4.2)   # Set to a given value
```


PromQL

Selecting series

Select latest sample for series with a given metric name:

```
node_cpu_seconds_total
```

 Open in PromLens

Select 5-minute range of samples for series with a given metric name:

```
node_cpu_seconds_total[5m]
```

 Open in PromLens

Only series with given label values:

```
node_cpu_seconds_total{cpu="0",mode="idle"}
```

 Open in PromLens

Complex label matchers (`=` : equality, `!=` : non-equality, `=~` : regex match, `!~` : negative regex match):

```
node_cpu_seconds_total{cpu!="0",mode=~"user|system"}
```

 Open in PromLens

<https://promlabs.com/promql-cheat-sheet/>

Aggregating over multiple series


Sum over all series:

```
sum(node_filesystem_size_bytes)
```

 Open in PromLens


Preserve the `instance` and `job` label dimensions:

```
sum by(job, instance) (node_filesystem_size_bytes)
```

 Open in PromLens

Aggregate away the `instance` and `job` label dimensions:

```
sum without(instance, job) (node_filesystem_size_bytes)
```

 Open in PromLens

Available aggregation operators: `sum()`, `min()`, `max()`, `avg()`, `stddev()`, `stdvar()`, `count()`, `count_values()`, `group()`, `bottomk()`, `topk()`, `quantile()`

PromQL sandbox

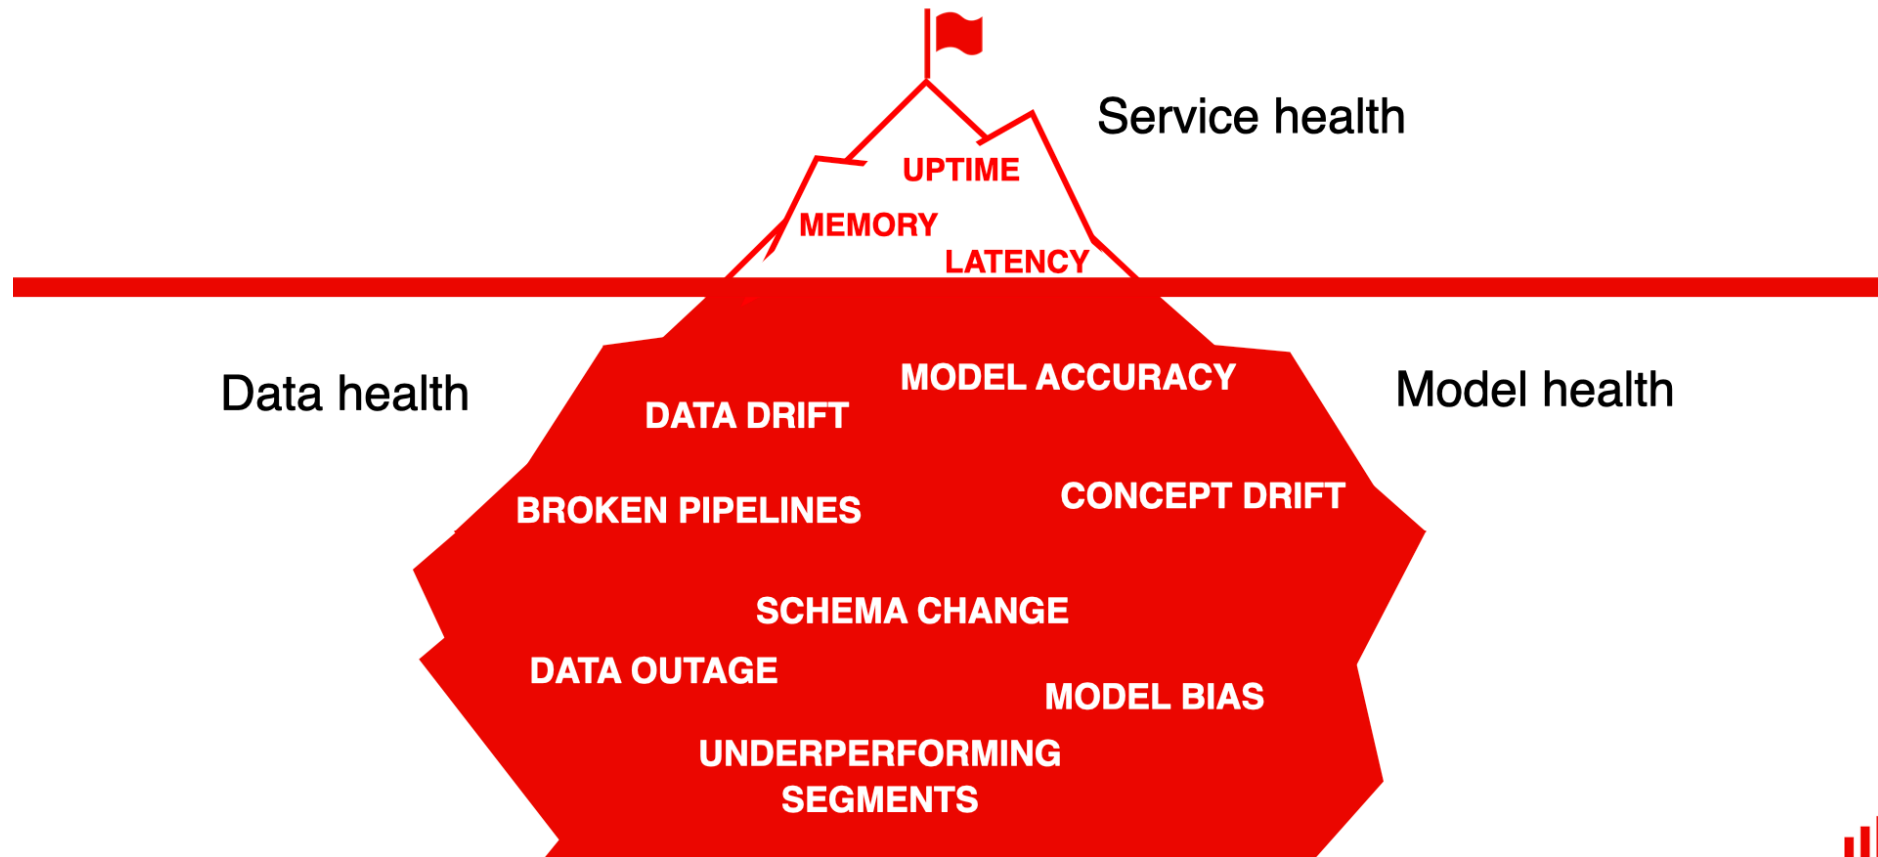
<https://demo.promlens.com>

Grafana



<https://play.grafana.org>

Что мониторить еще?



Что мониторить еще?

Прекрасные референсы:

<https://www.evidentlyai.com/ml-in-production/data-drift>

<https://www.evidentlyai.com/ml-in-production/model-monitoring>

<https://www.evidentlyai.com/ml-in-production/concept-drift>

Data Drift

Incoming data



Feature distribution: sales_channel



Model quality: accuracy over time



Concept Drift

Data drift

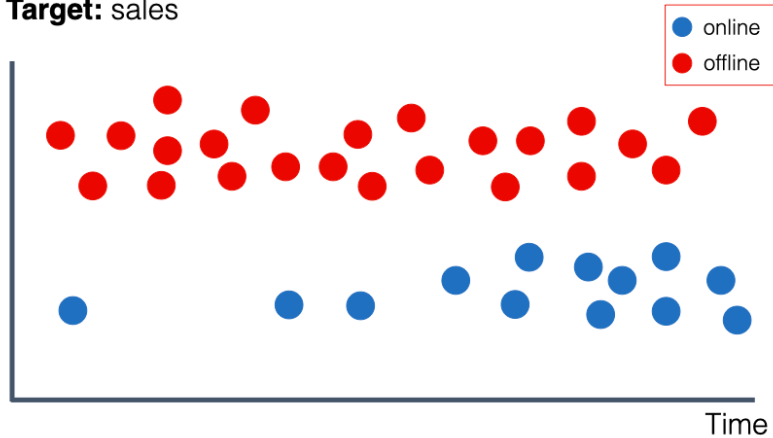
Concept drift

$P(X)$

$P(Y|X)$

Data drift

Target: sales

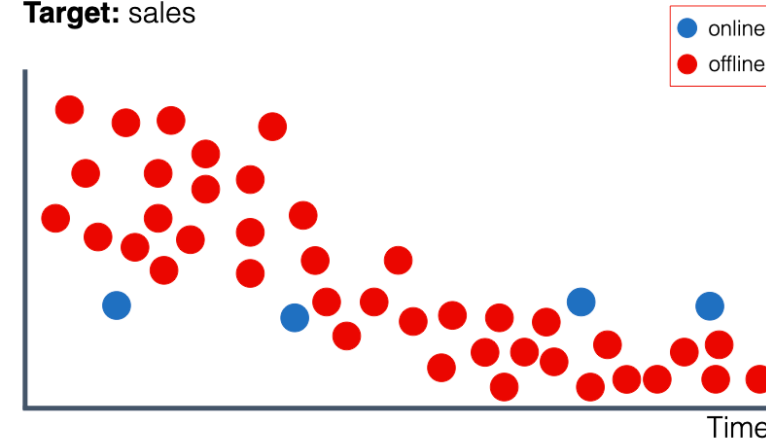


Feature distribution: sales channel



Concept drift

Target: sales



Feature distribution: sales channel



Data drift

a	a	a	a	a	n	a	a	M
b	b	b	b	b	b	b	b	b
c	n	c	c	c	c	n	c	m
d	d	d	d	d	d	d	N	d
e	e	e	m	e	e	h	N	e
f	f	f	f	f	f	f	f	n

Time

Train-serving skew


a	a	a	a
b	b	b	b
c	c	c	c
d	d	d	d
e	e	e	e
f	f	f	f

Training


a	a	A	a
l	B	b	l
c	C	C	L
d	d	d	d
e	e	e	e
f	F	F	P

Production


Data drift




Data quality



Data drift



Outliers





Projects / Sales forecasting

Last 14 days

Daily predictions



Share of drifting features



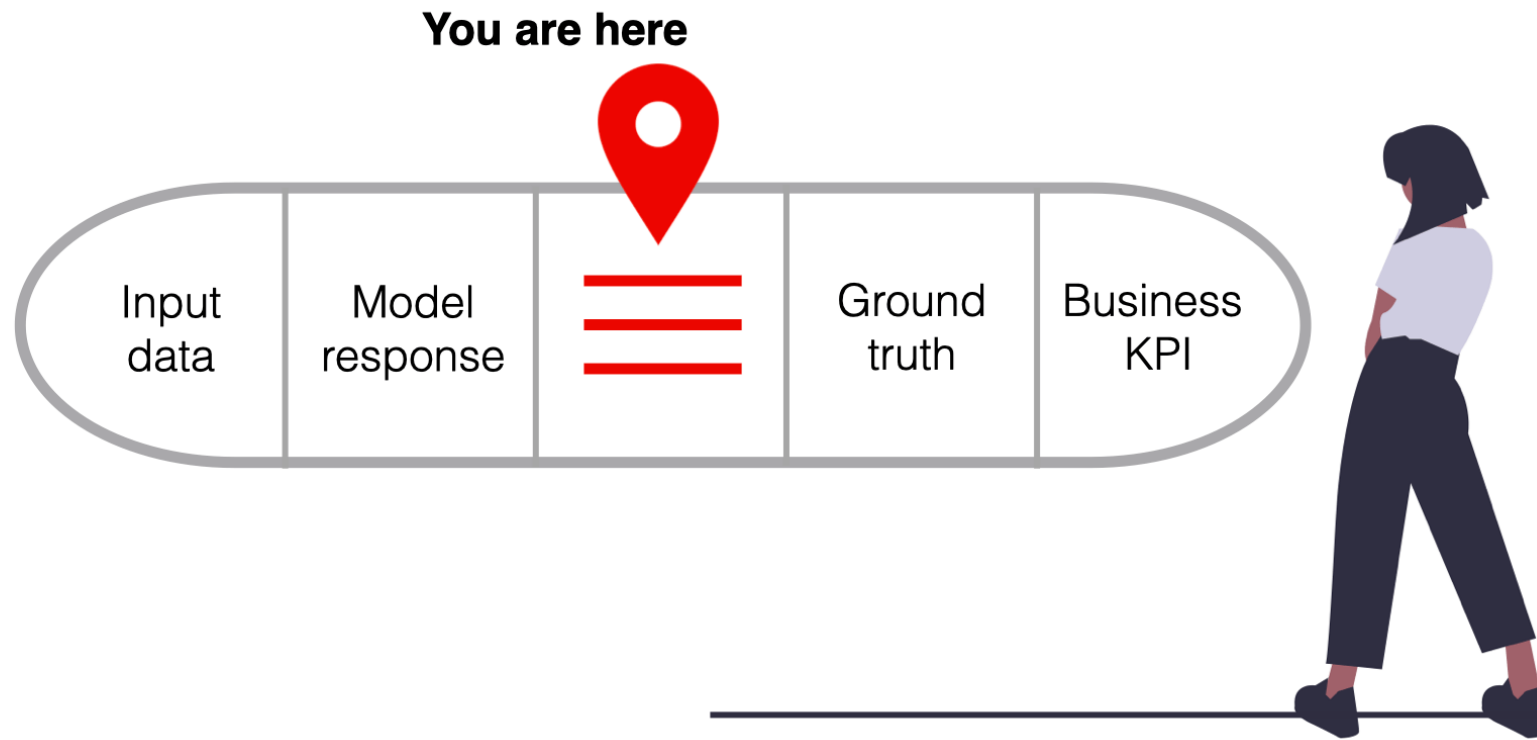
Prediction drift



Null values %




GT delay




GT delay

Accuracy over time




Share of **drifting** features



Data Drift Detect. Summary Statistics


Ship_Mode
cat

	current	reference
count	0	976
unique	0 (0.0%)	4 (0.2%)
most common	nan (100.0%)	nan (51.2%)
missing	2000 (100.0%)	1024 (51.2%)
new categories	0	
missing categories	4	

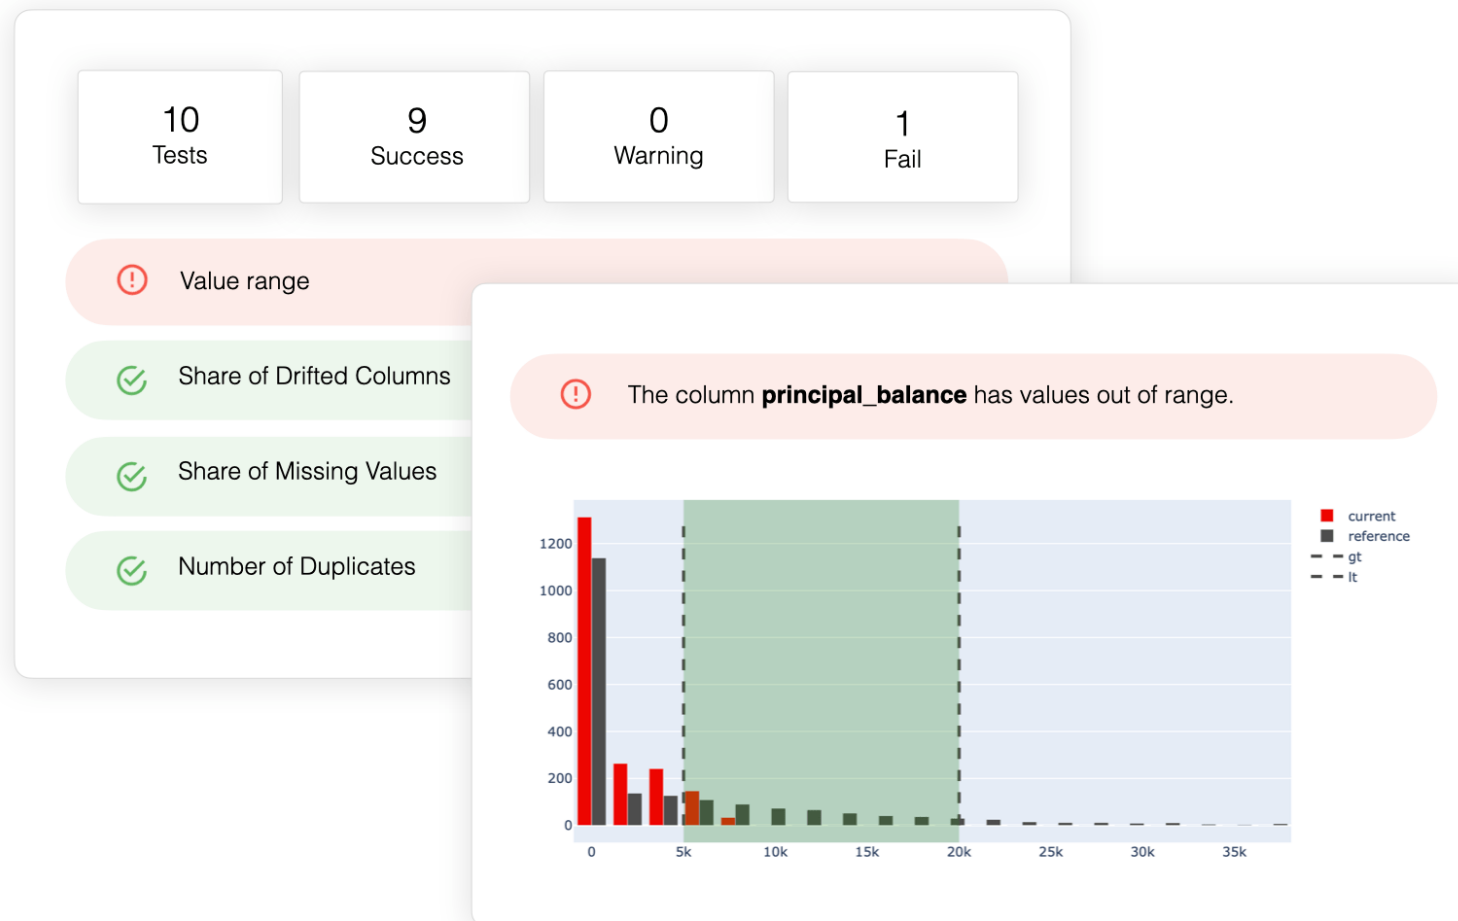


State
cat

	current	reference
count	2000	2000
unique	3 (0.15%)	48 (2.4%)
most common	California (60.2%)	California (19.95%)
missing	0 (0.0%)	0 (0.0%)
new categories	0	
missing categories	45	



Data Drift Detect. Summary Statistics




Data Drift Detect. Статистические тесты



Few important features

VS



A lot of unimportant features



Data Drift Detect. Distance metrics

Drift in 'amount_borrowed'
Data drift detected.
Drift score: 1.01 (Wasserstein distance).

DATA DRIFT

DATA DISTRIBUTION

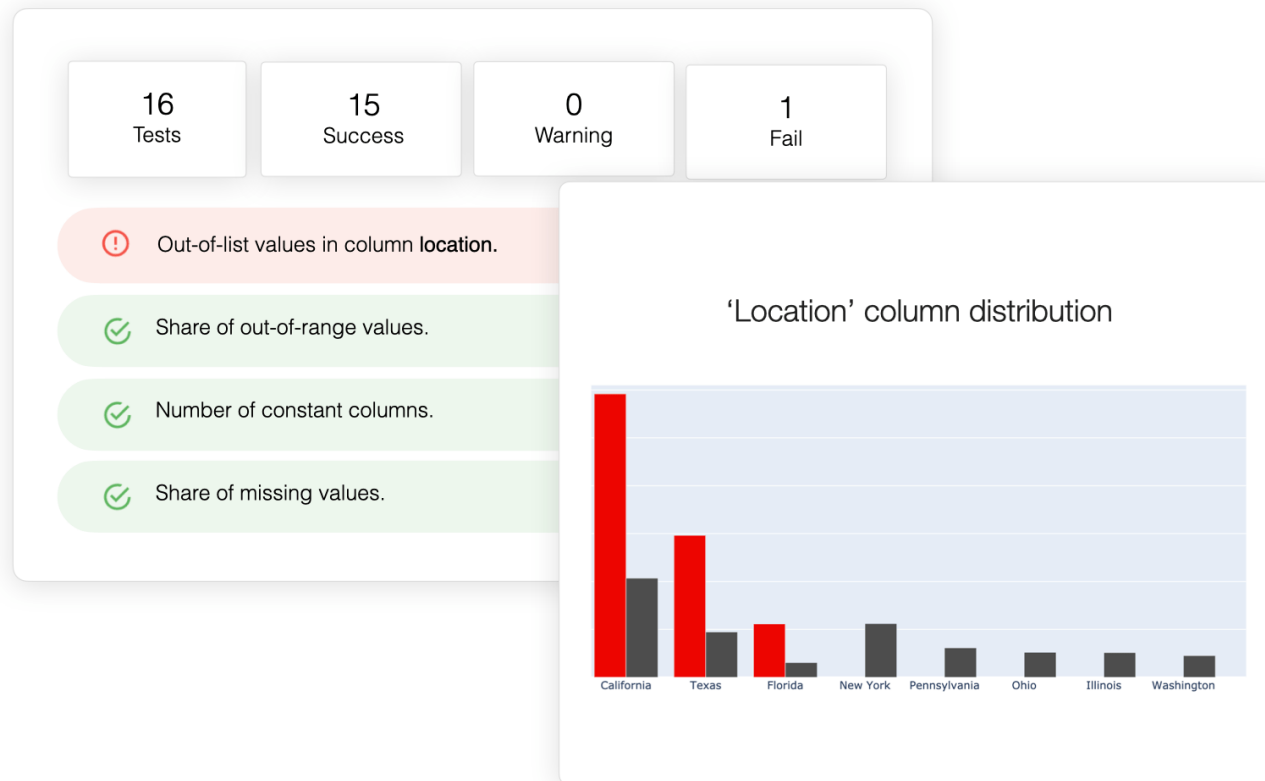


Dataset drift



Drift is detected for 40.0% of columns (6 out of 15).

Search						
Column	Type	Reference Distribution	Current Distribution	Data Drift	Stat Test	Drift Score
> class	cat			Detected	Jensen-Shannon distance	0.273594
> purchase_value	num			Detected	Wasserstein distance (normed)	0.389834
> source	cat			Detected	Jensen-Shannon distance	0.296321
> ip_address	num			Not Detected	Wasserstein distance (normed)	0.048344
> user_id	num			Not Detected	Wasserstein distance (normed)	0.036559
> browser	cat			Not Detected	Jensen-Shannon distance	0.030073

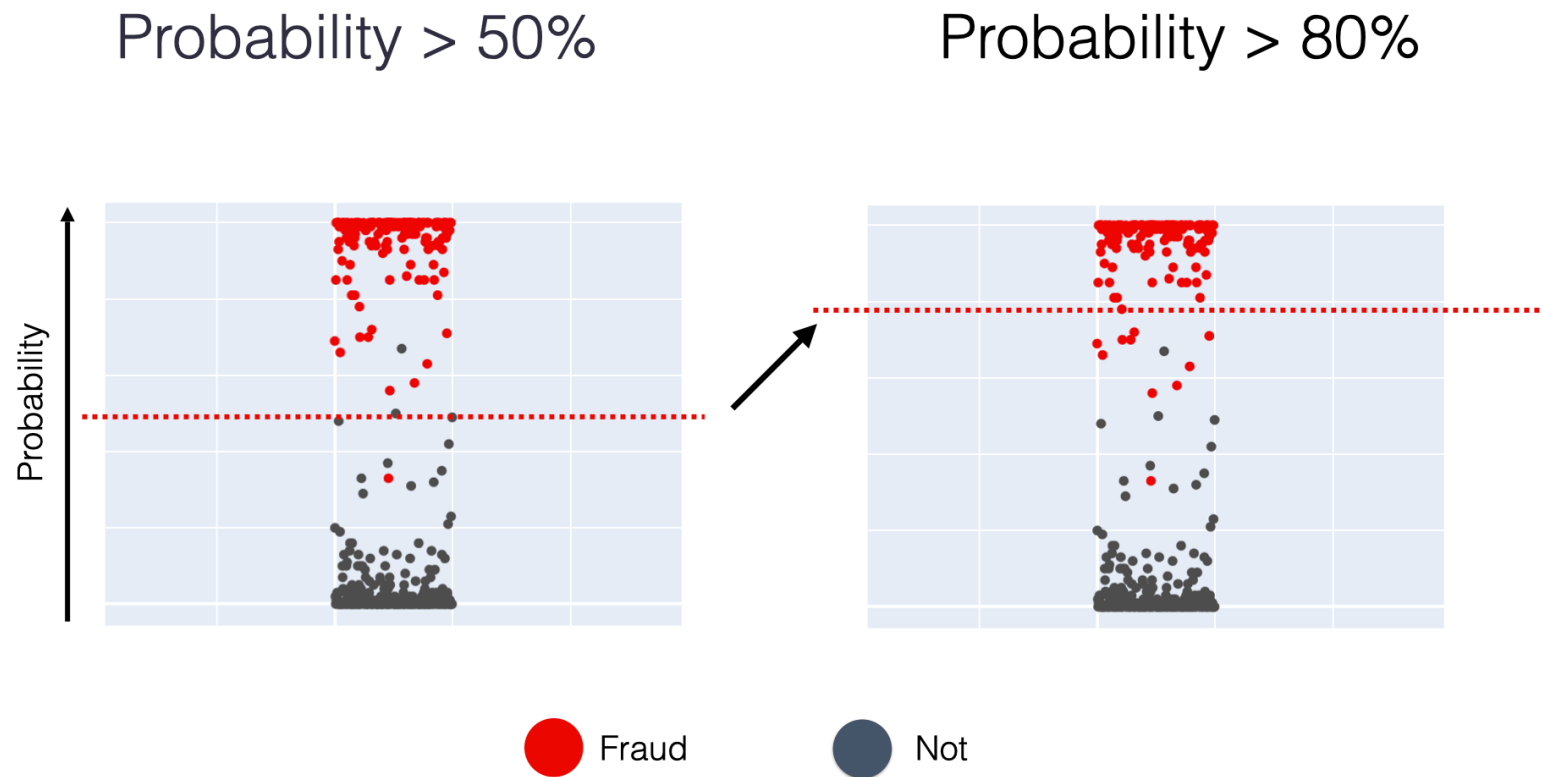
Data Drift Detect. Rule-based checks



Retraining



Process intervention



Evidently Demo

<https://demo.evidentlyai.com>